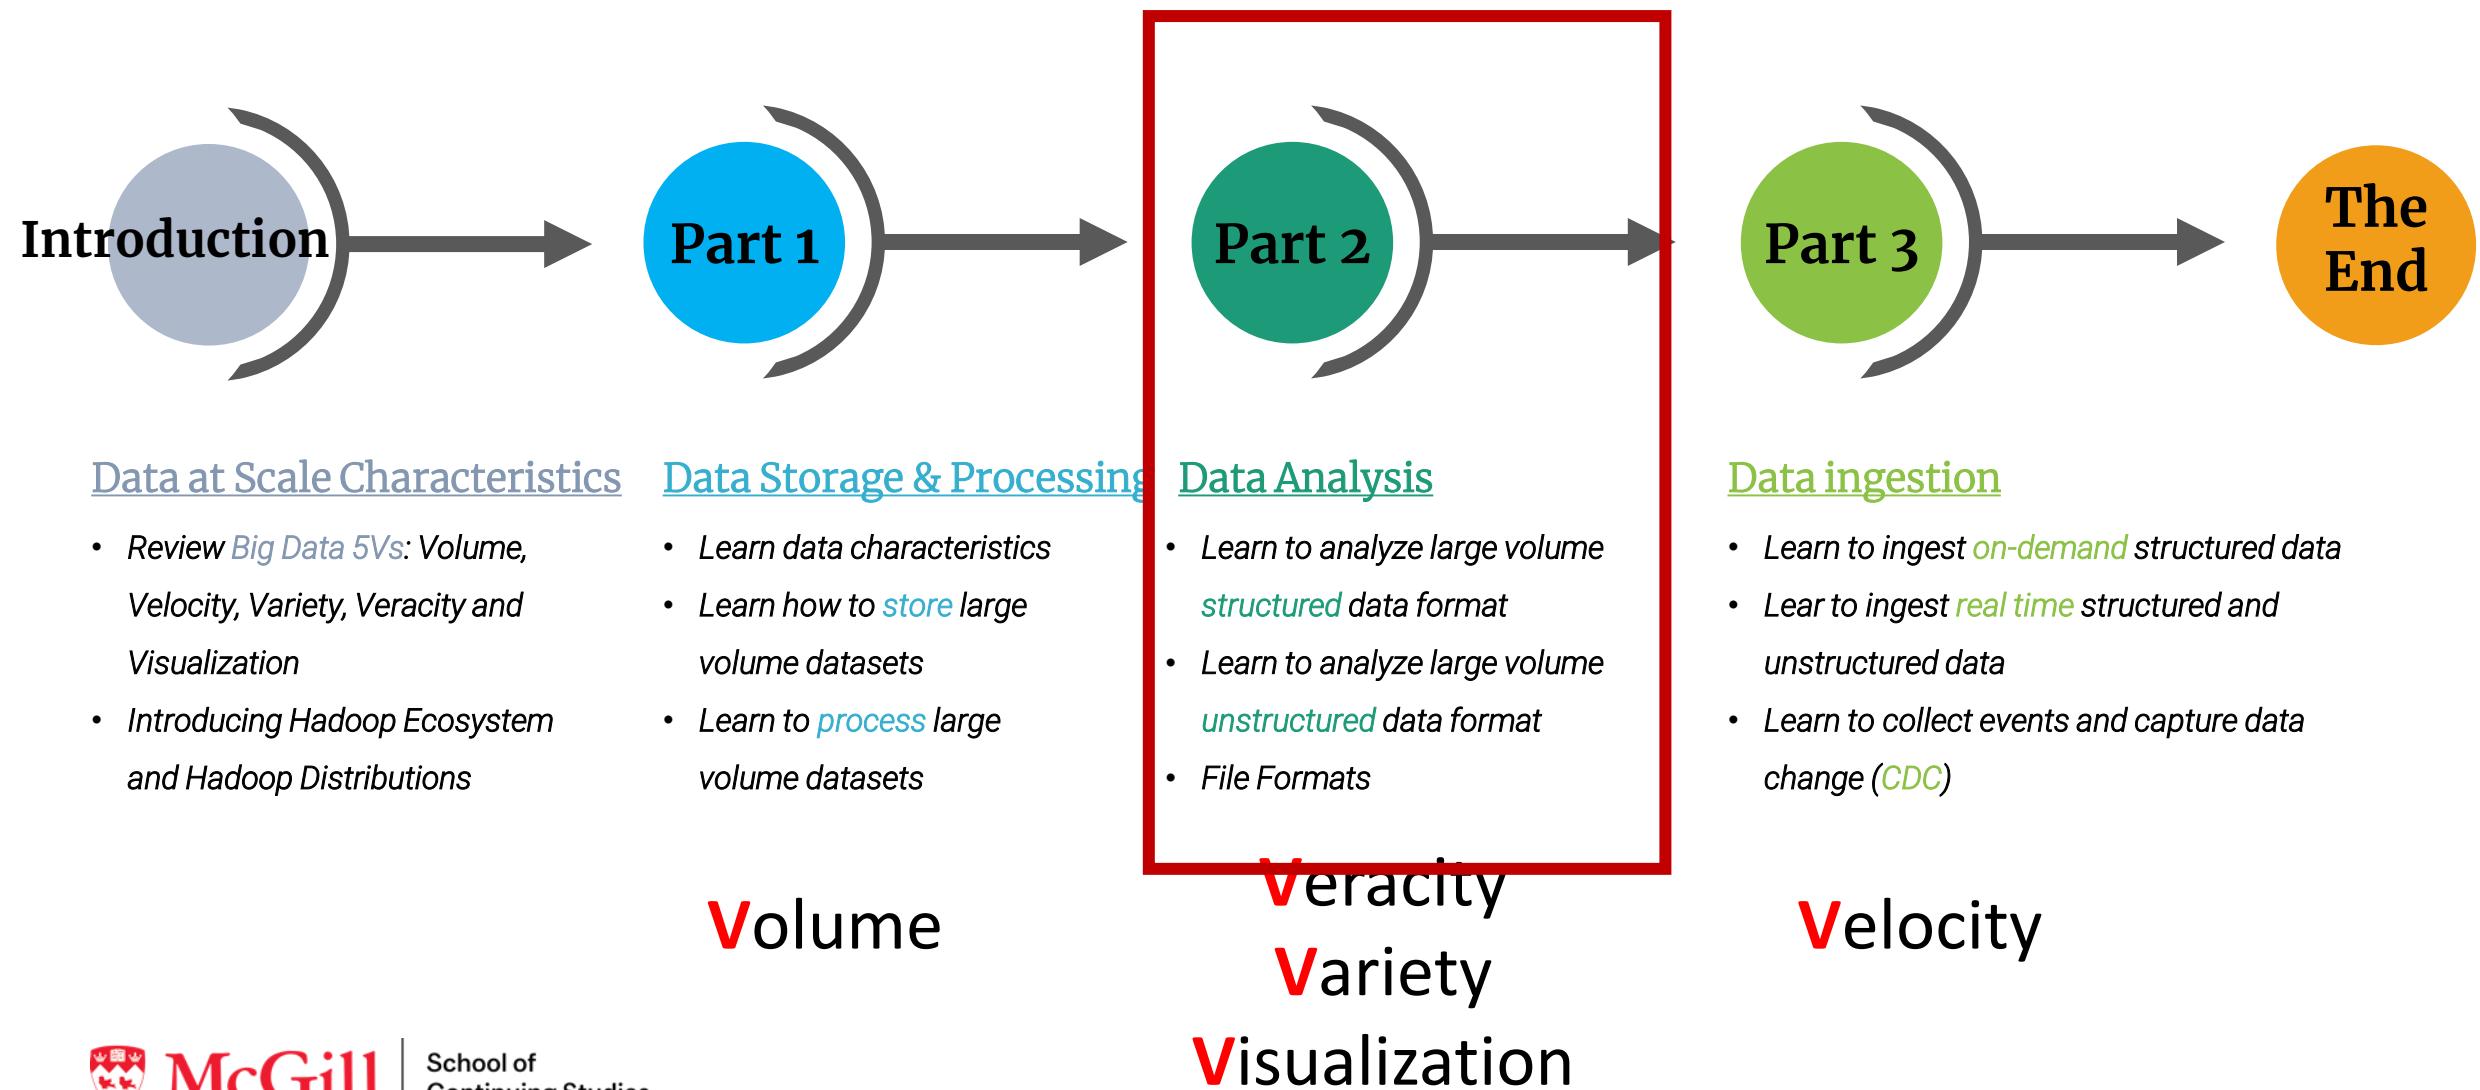


YCBS 257 – Data at Scale

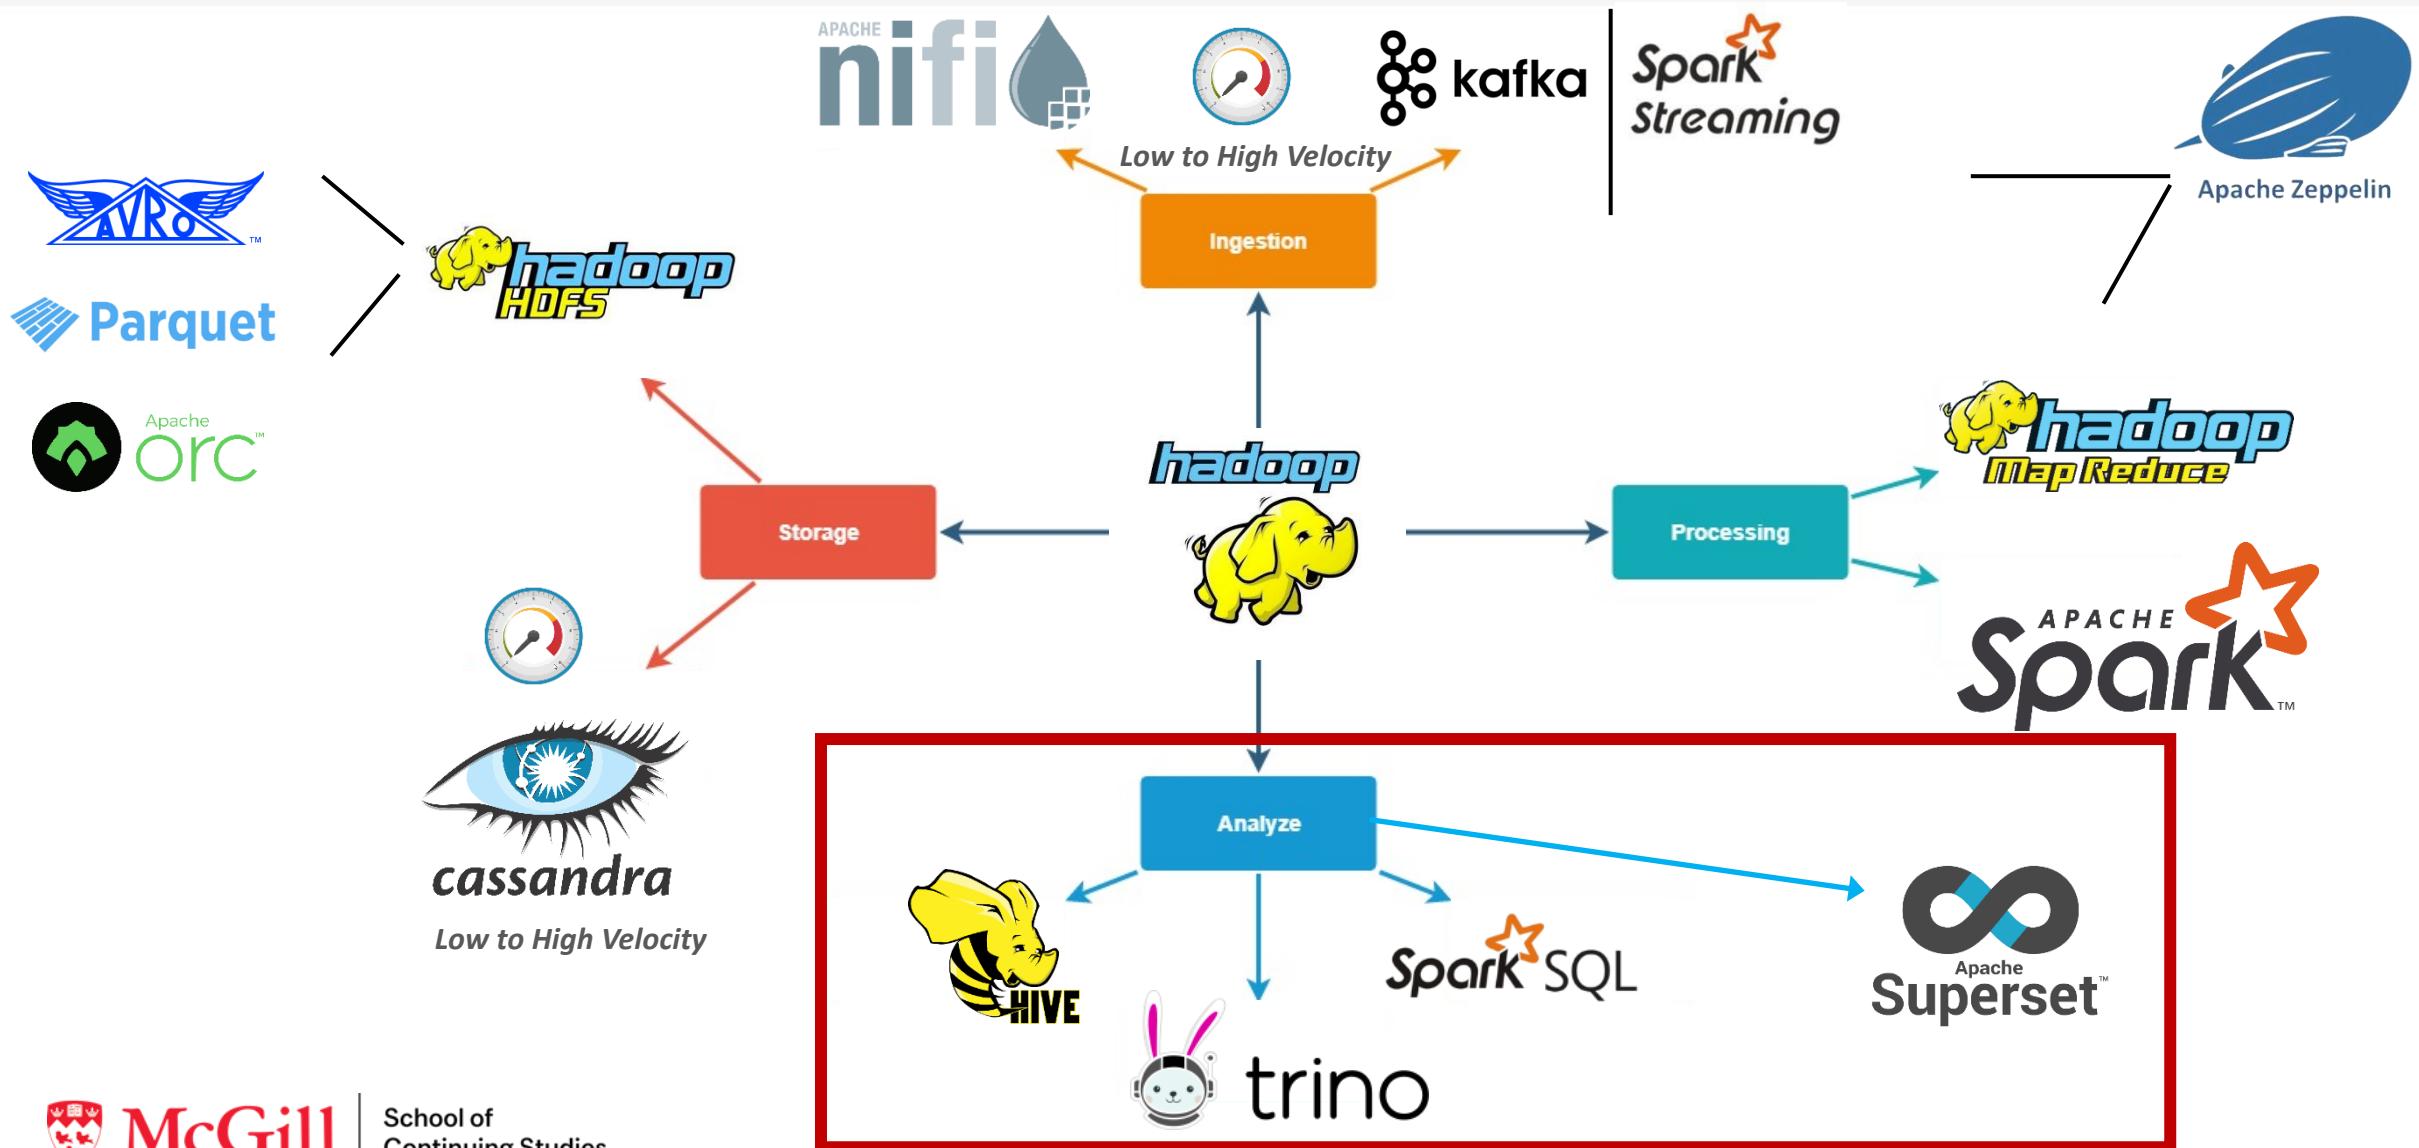
Instructor: **Dr Khaled Tannir**



Course Learning Path (CLP)



Data at Scale Popular Technologies

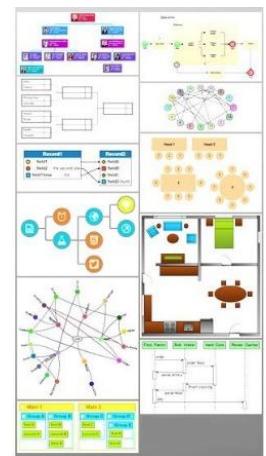
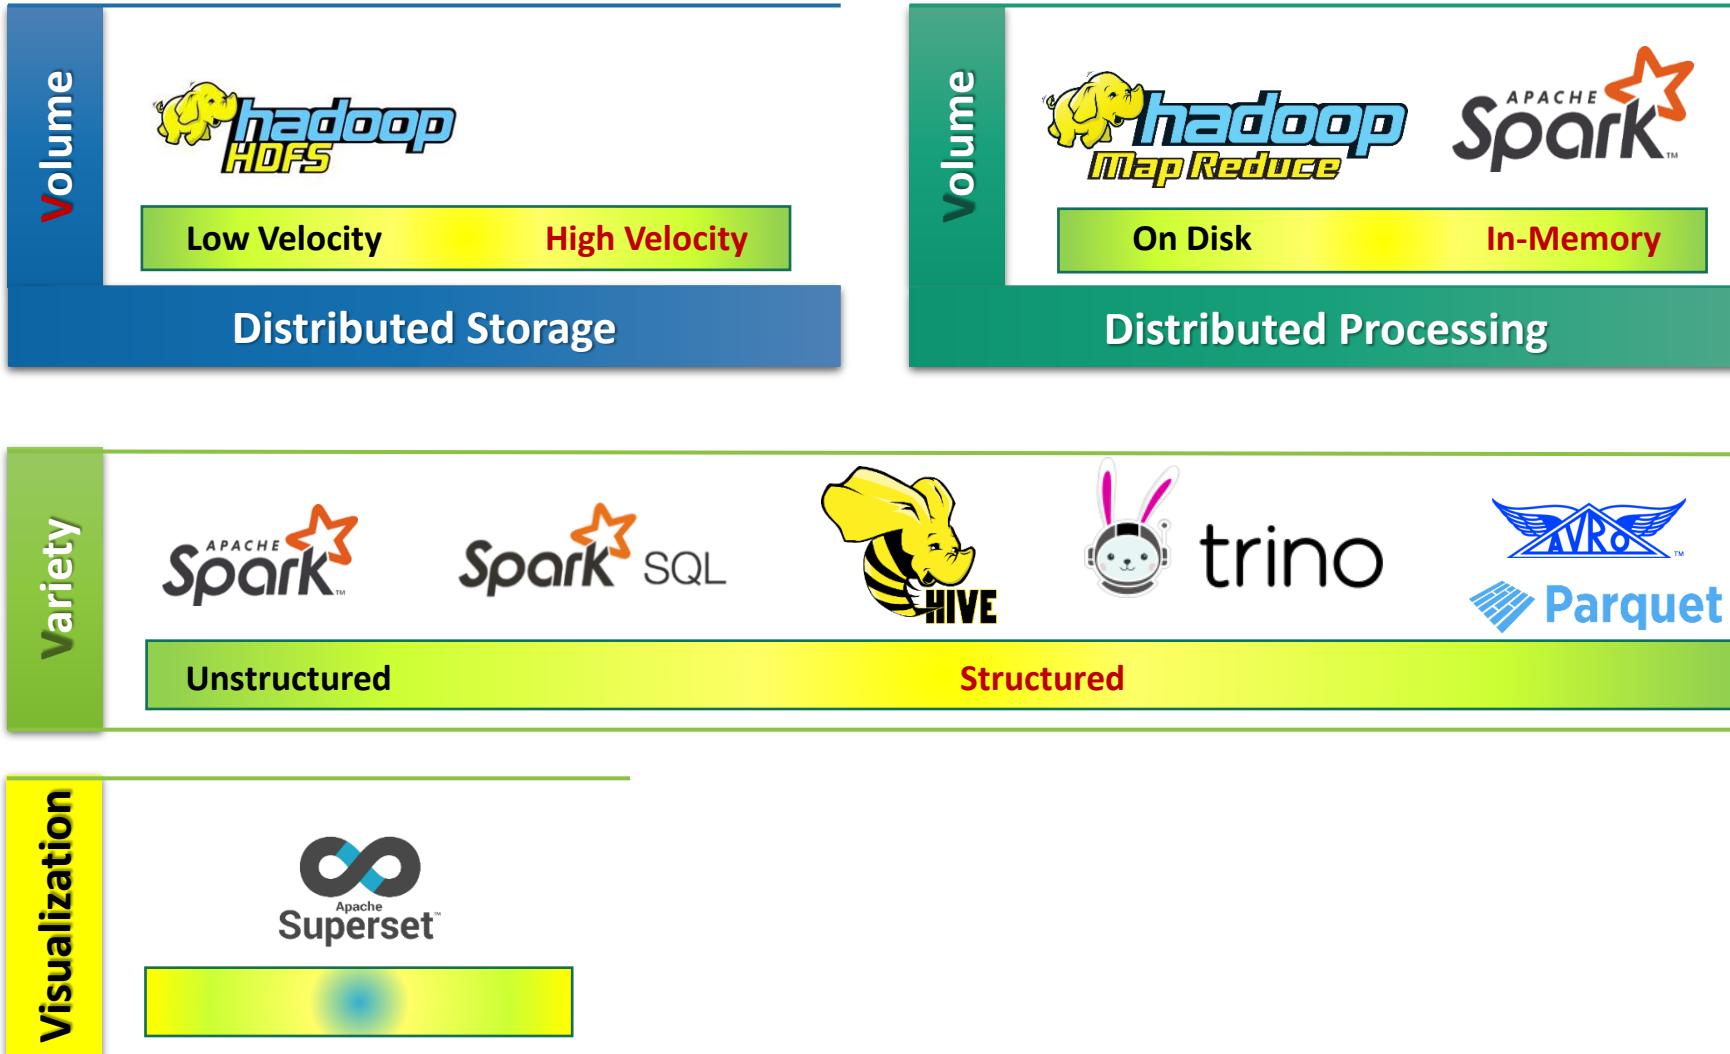


Data at Scale Course

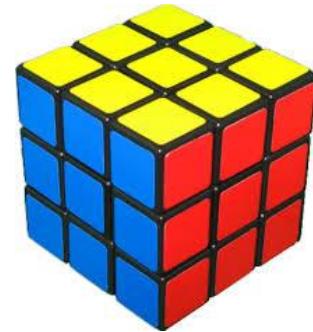
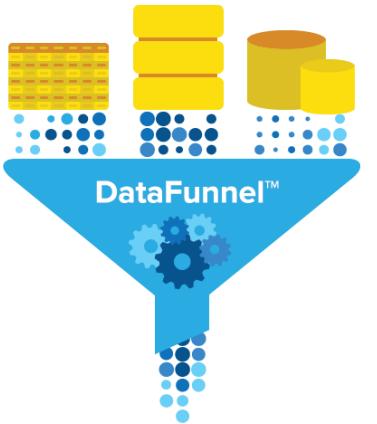
Part 1



Part 2



The Big Data Process Model



- yes
- no
- maybe

1 Collect

Collect data from sources

- Open data
- Traditional ETL
- ERP / Data warehouse
- RDBMS
- ...

2 Organize

Prepare the data

- Cleansing
- Filtering
- Sorting
- Define the format
- Define the lifecycle.
-

3 Analysis

Extract Knowledge from the data

- Find insights
- Modeling behavior
- Probabilistic rather than definitive.
- Text, voice, image analysis
- ...

4 Decide

Use insights to change business

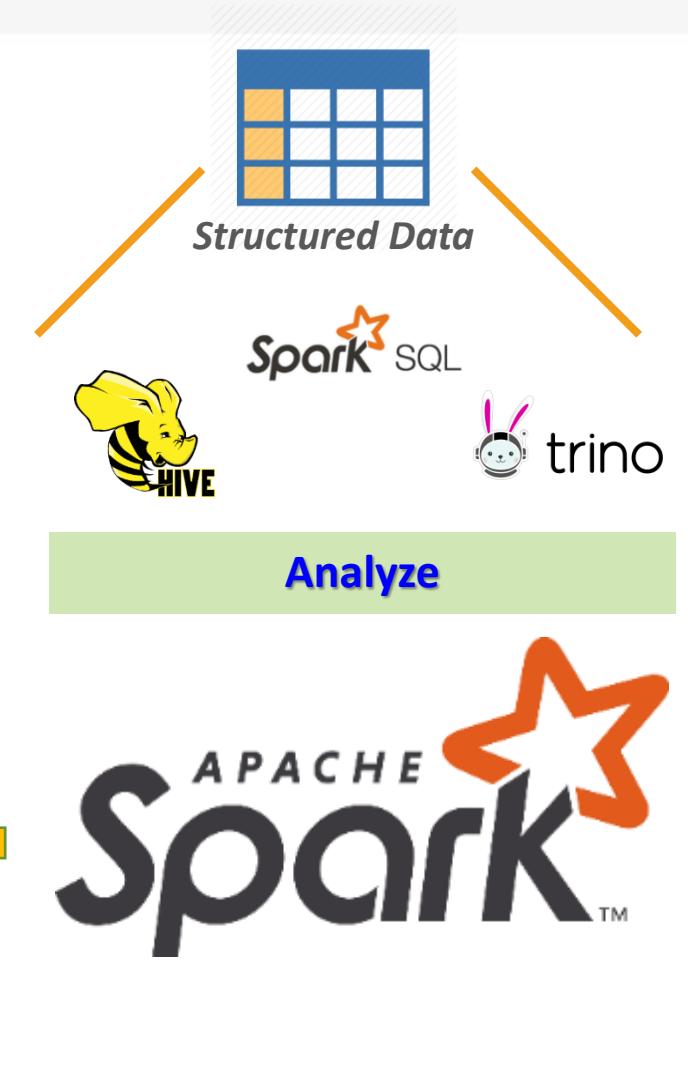
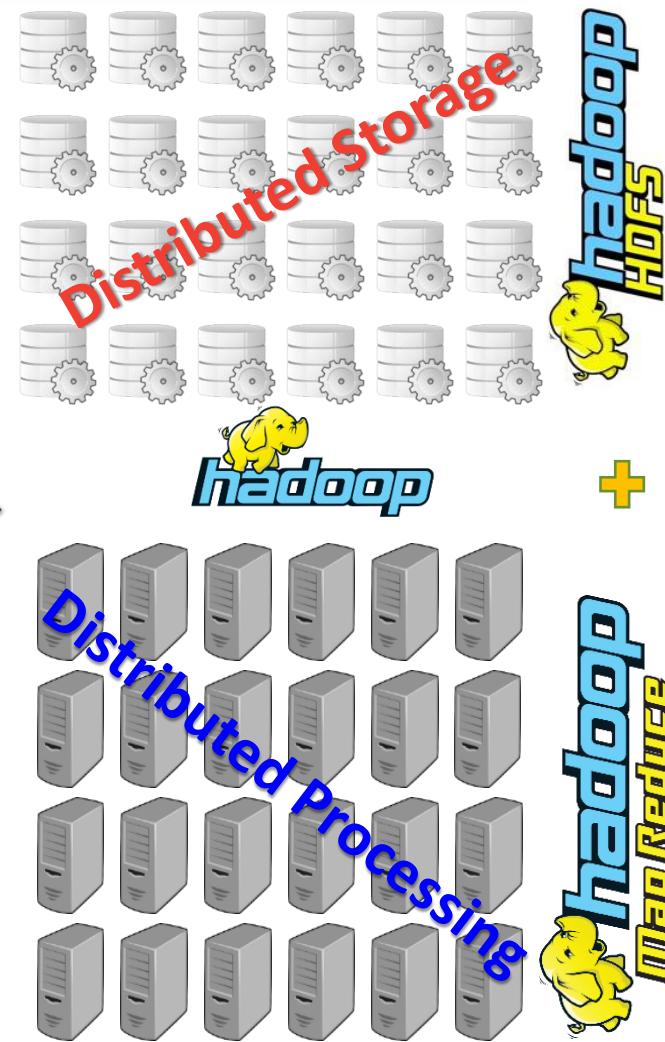
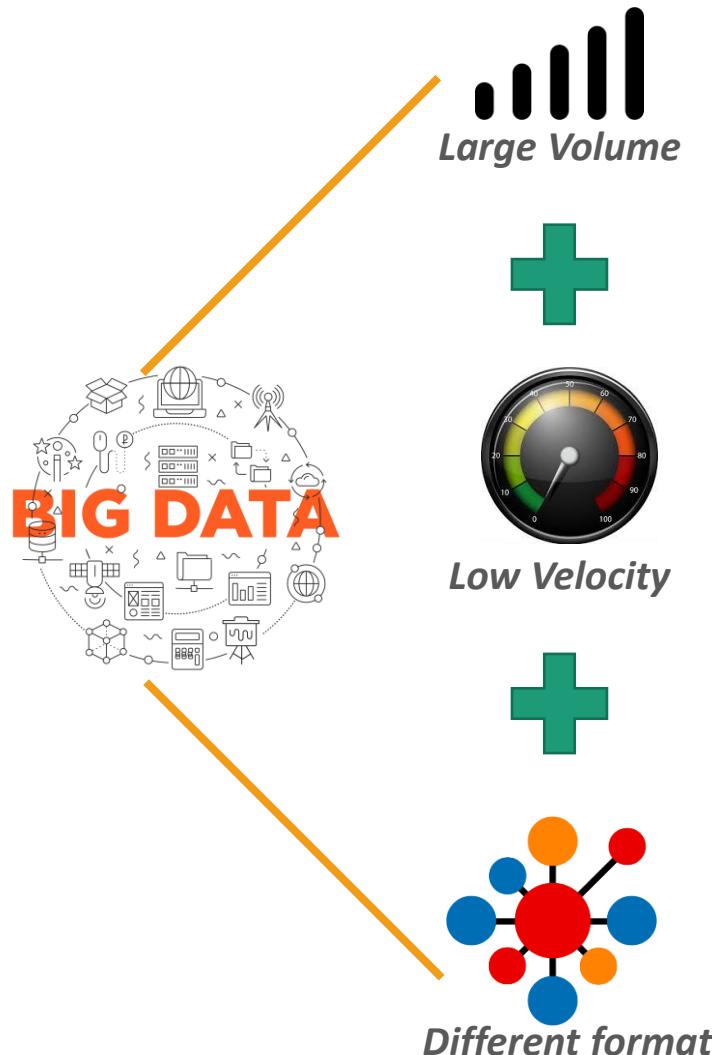
- Reports
- Dashboards
- BPM – Real time decision
- Information sharing
- ...

Theme of this Course

- ***Analyzing structured Data at Scale Concepts***
- ***Apache HIVE***
 - ***Core Concepts***
 - ***Hive Architecture***
 - ***Hive Data Model***
 - ***HiveQL query Language Essentials***



Data at Scale



Popular Tools to Analyze Data at Scale



FRAMEWORKS



FORMAT



QUERY / DATA FLOW



DATA ACCESS



DATABASES



OLAP



MAD Landscape 2023 - <https://mattturck.com/mad2023/>

Analyzing Data at Scale



- ***How to analyze a distributed file / Dataset ?***
How to organize/query data in a distributed environment?
- ***What if you have to analyze billions of rows every day?***

How long will take analyzing your dataset?

Analyzing Data at Scale



- To analyze/query large / distributed datasets you need to consider

1

Data format (ORC, Avro, Parquet, ...)

Reduce the volume and Optimize the structure of datasets

2

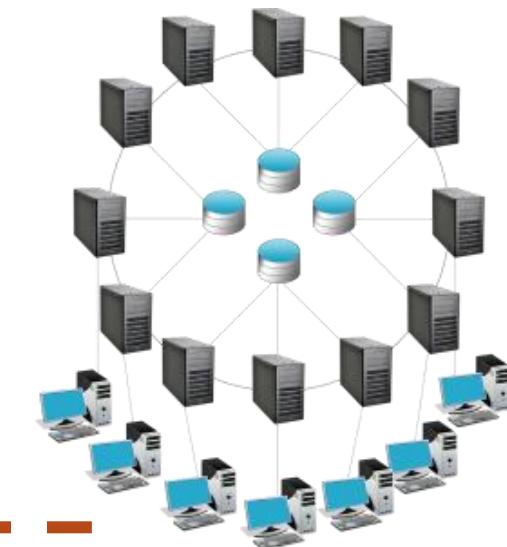
Data Partitioning

Decomposing datasets into more manageable parts

3

Data Bucketing (Hive/Spark)

Provides a way of segregating equally data into multiple partitions (files/directories)





1- Data Format

- CSV, JSON, XML, TXT, ... Human readable, doesn't mean that's the best way to store data (*High disk space usage*)
- Those file formats cannot be stored in a parallel manner
(*Cannot parallelize data processing*)
- Avro, ORC, and Parquet Three optimized file formats (*ORC=Optimized Row Columnar*)
- Each of them is unique and bring its own relative advantages and disadvantages

FORMAT

ICEBERG

Parquet

HUDI

ORC

ARROW

DELTA LAKE

AVRO



2- Data Partitioning

- HDFS does not have the concept of indexes.
- Even for reading one row, the entire file should be read.

High disk I/O and delays

- **Partitioning** provides a way to read only a subset of data.

Based on a Partition key (Column Name)

- Multiple attributes can be used for hierarchical partitioning.
- Split data into directories based on individual values of attributes.

Partition value (Column Value)



3- Data Bucketing

- Partitioning is **only optimal when a given attribute has a small setup unique values**

What if we need to partition for a key with a large number of values without proliferating the number of directories?

- **Bucketing** is similar to Partitioning (Even distribution)
- It uses a hash function instead of the value of the attribute to convert the value into a special hash key

Values that have the same hash key go to the same bucket

- Number of Buckets can be controlled
- Ideal for attributes with large number of unique values

What is Apache Hive?



- A Data Warehouse on top of Hadoop (**SQL on Hadoop**).
- Can deal with different **storage** and **file** formats using ***Input/Output*** and ***SerDes***.
- Summarize Big Data and makes querying and analyzing easy.
- Familiar, scalable, and extensible.
- Written in Java and open-source



Features of HIVE



- Stores **schemas** in a database (metastore) and processes data stored in HDFS.
- Supports **User-Defined Functions (UDF)**, scripts, and a customized I/O format to extend its functionality.
- Supports running on different computing frameworks (*MapReduce, Tez, Spark..*)
- Provides **HiveQL** (SQL like Query Language) and Rich data types (structs, lists and maps).

Efficient implementations of SQL filters, joins and group-by's on top of MapReduce

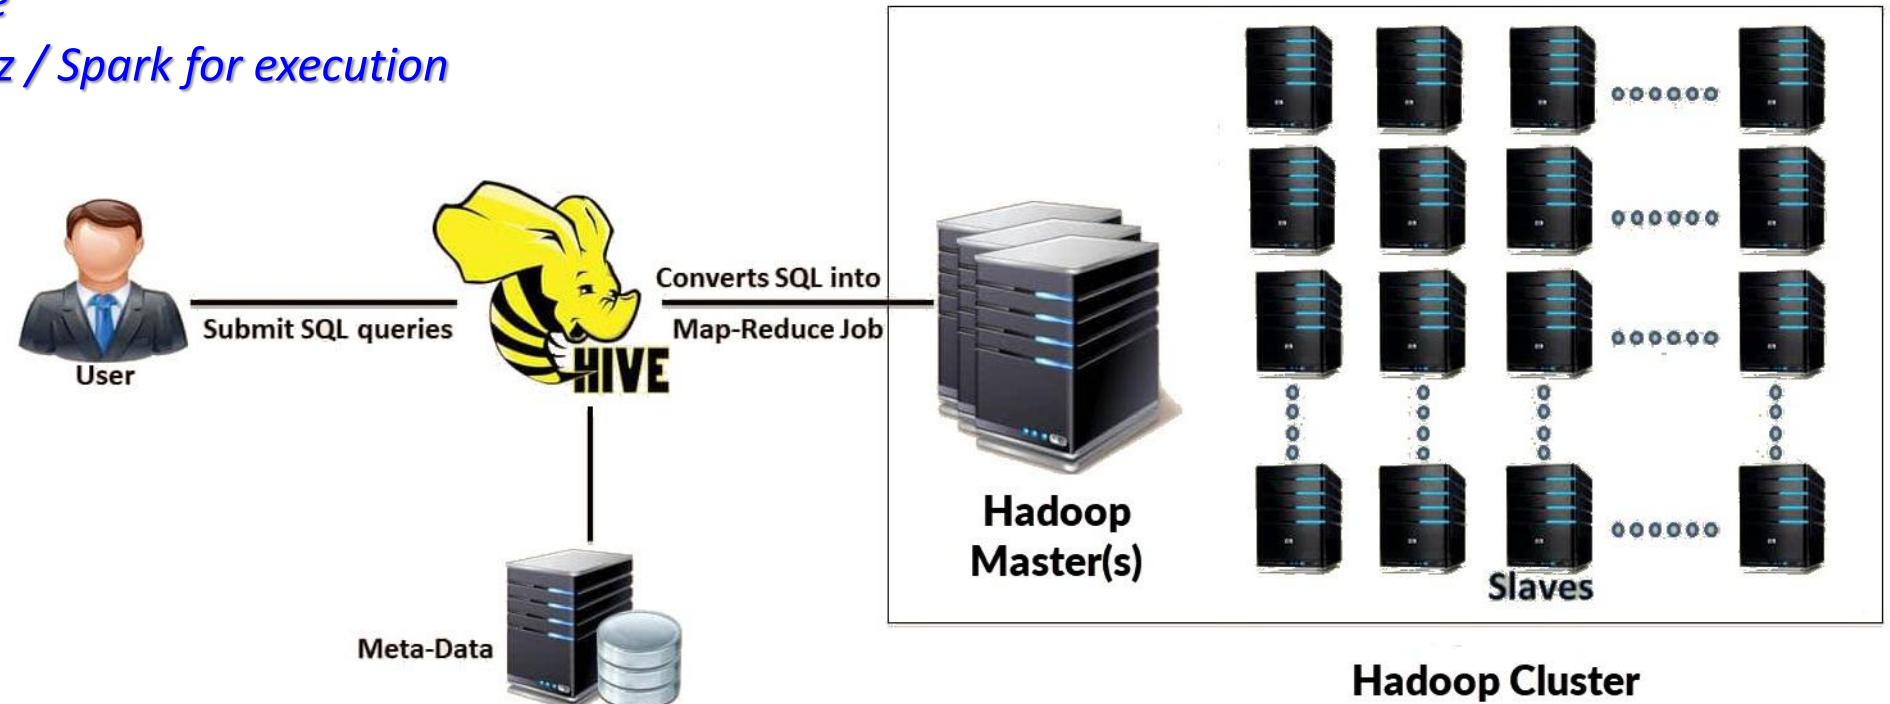
- Hive is able to query PB of data.



How it Works?

● Hive is built on top of Hadoop

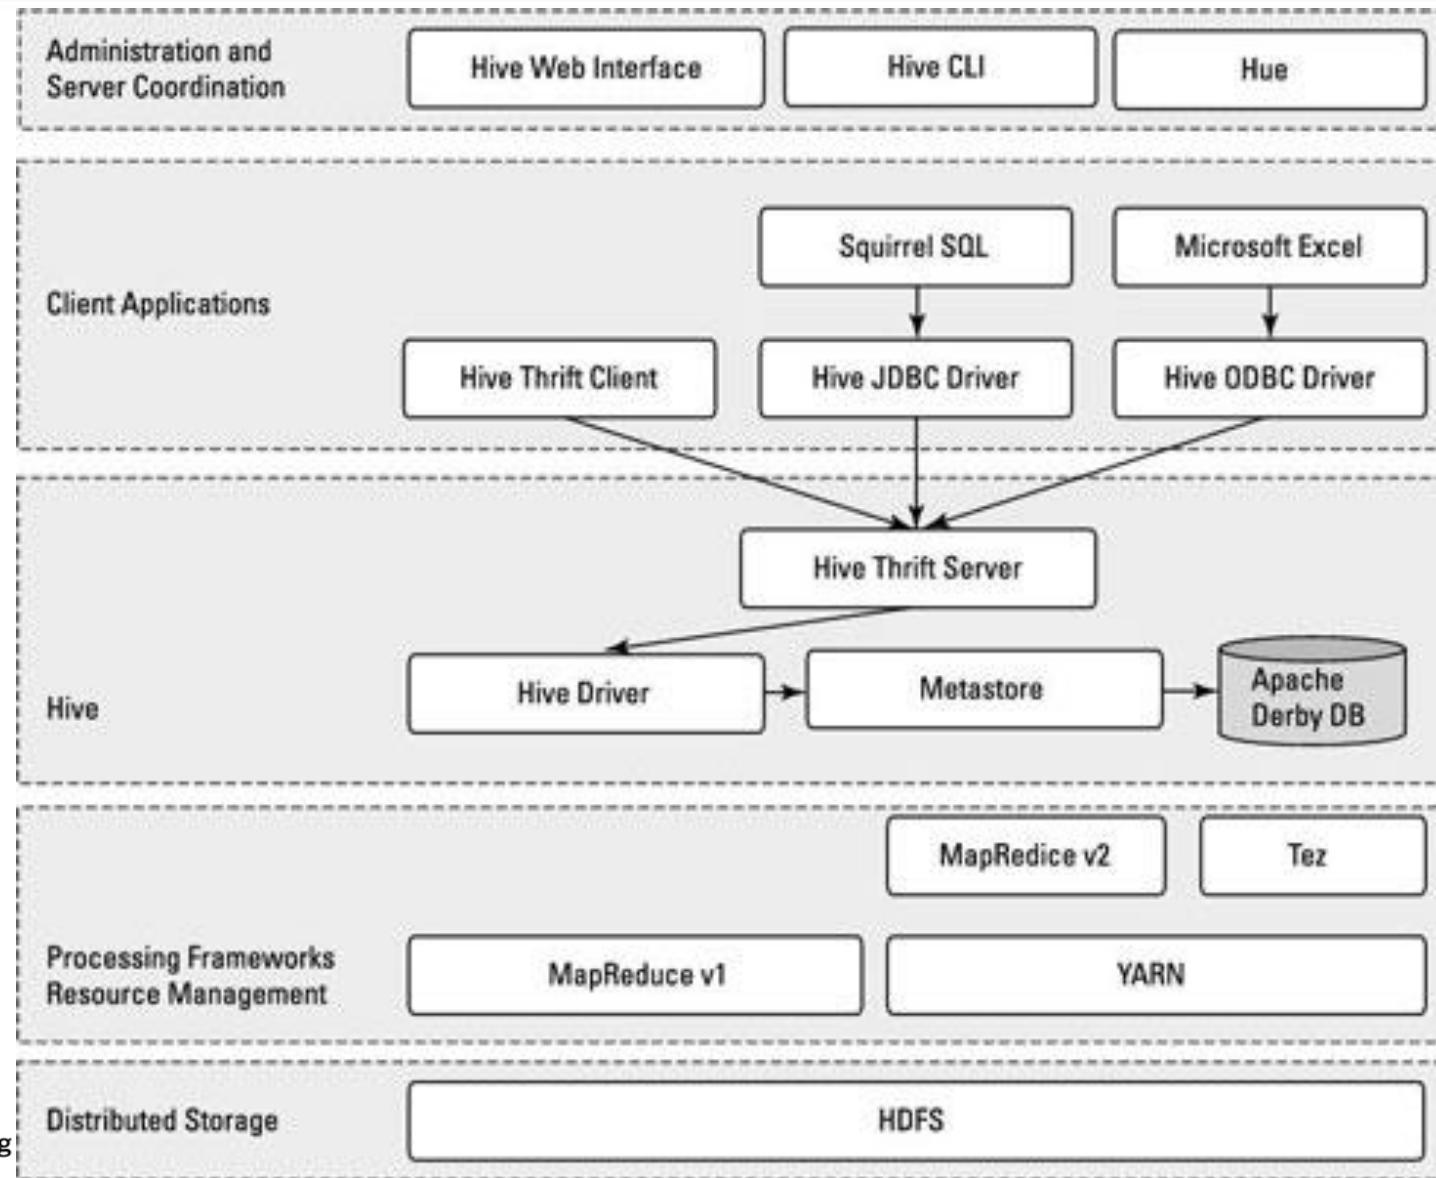
- *Uses HDFS for storage*
- *Uses MapReduce / Tez / Spark for execution*



● Hive compile HiveQL queries into MapReduce / Tez / Spark jobs and run the jobs in the Hadoop cluster

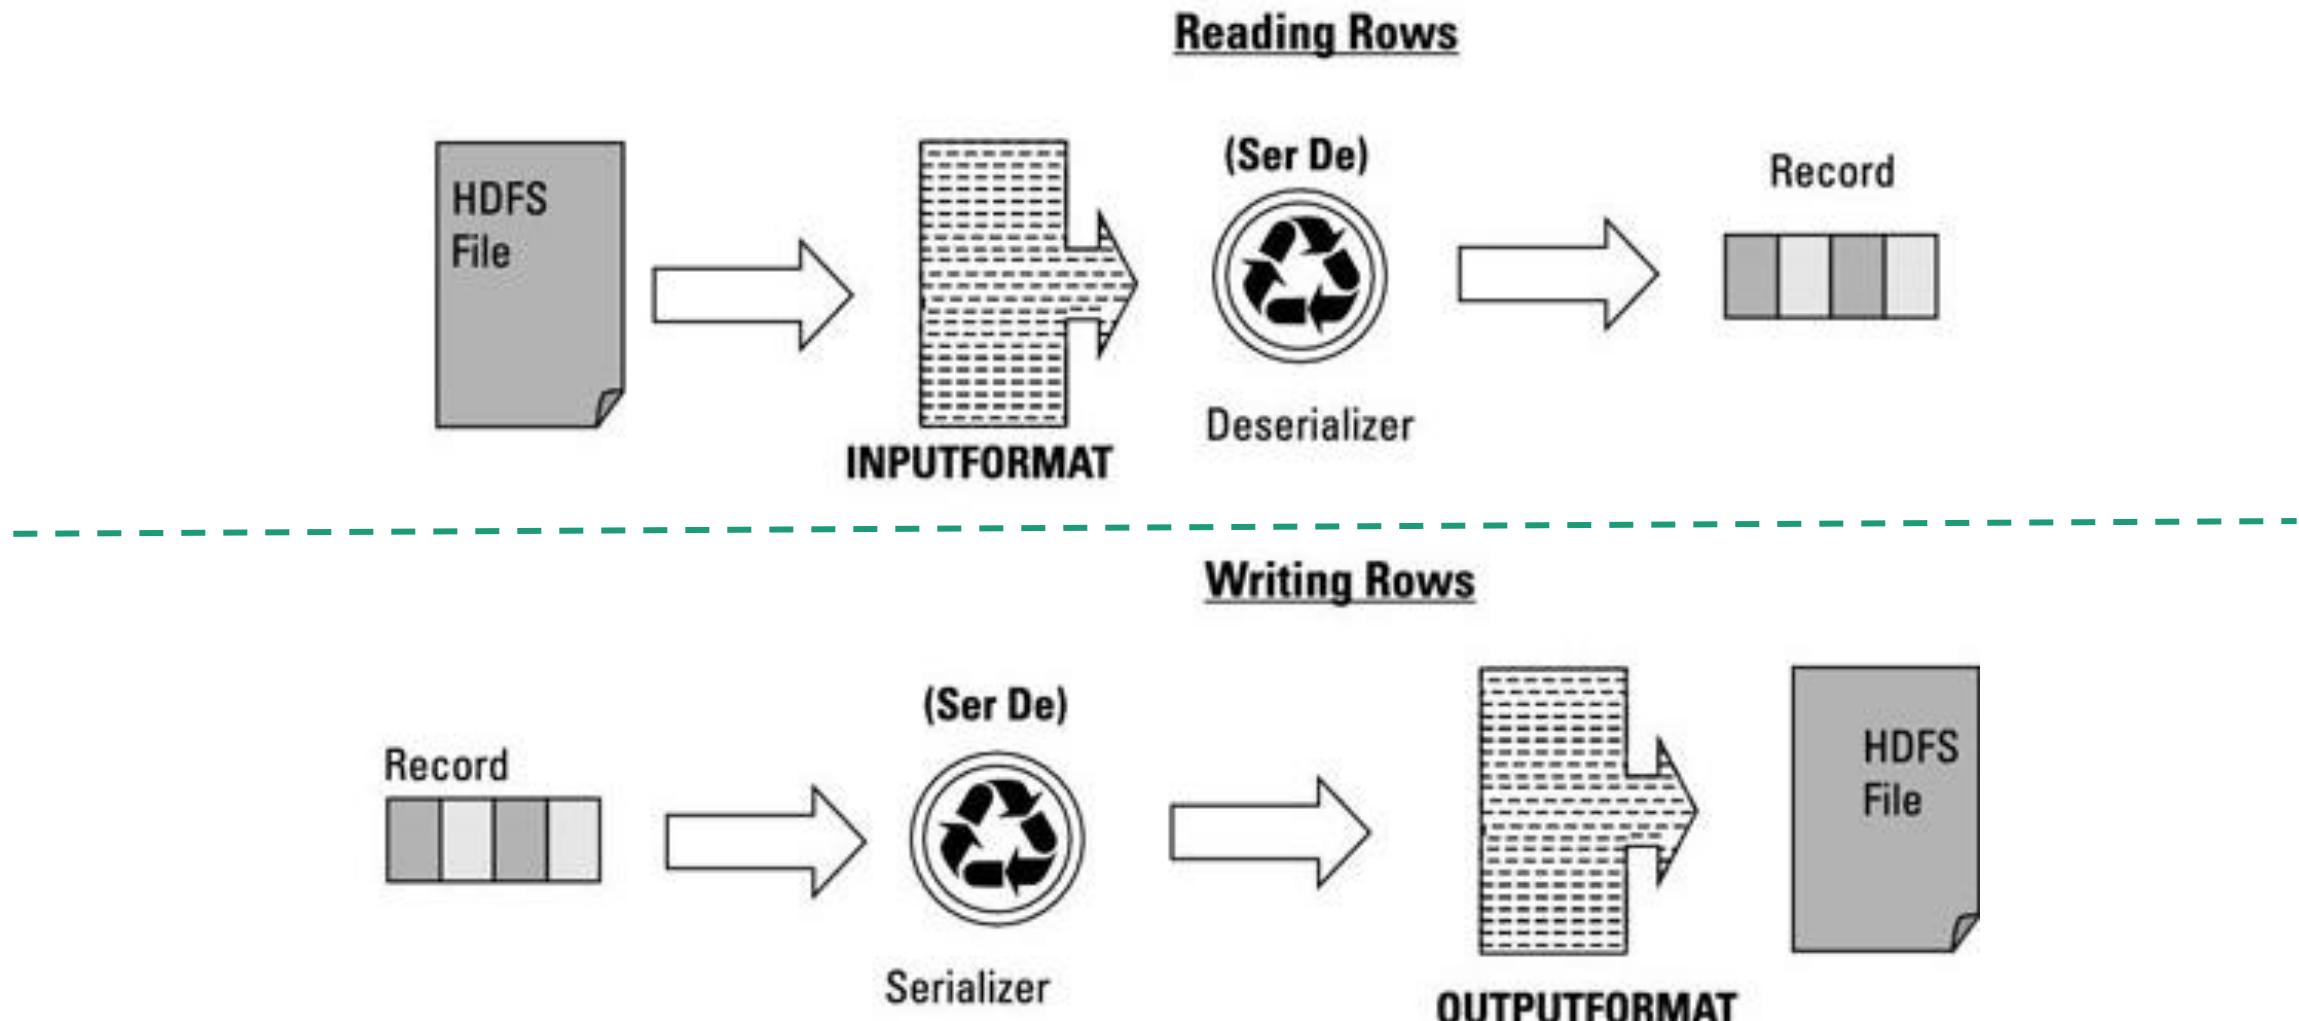


Hive High Level Architecture





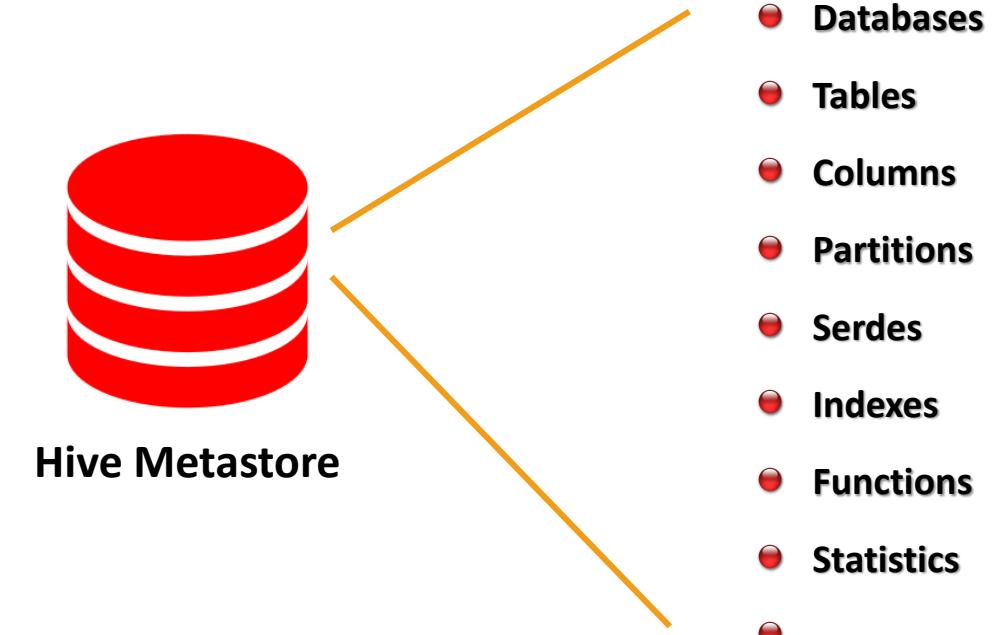
Hive Read/Write Operations



Hive Metastore



- Internally stored in a relational database
- A service that provides metastore access to other Apache Hive services
- Disk storage for the Hive metadata which is separate from HDFS storage.





Hive Data Model

Database

Namespace containing a set of tables

Tables

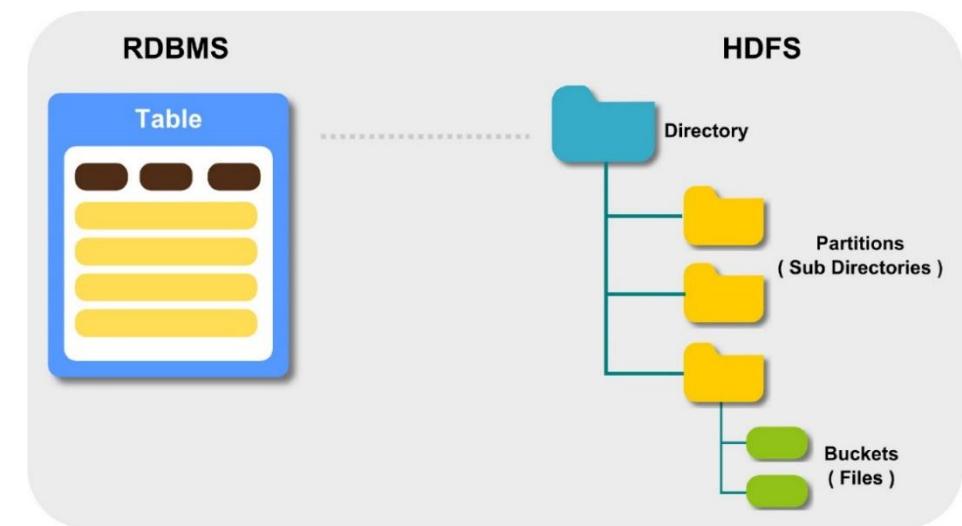
- *Analogous to relational tables*
- *Each table has a corresponding directory in HDFS*
- *Data serialized and stored as files within that directory*

Partitions

- *Each table can be broken into partitions*
- *Stored in a sub-directory within a table's directory*

Buckets

- *Data in each partition divided into buckets*
- *Based on a hash function of the column*
- *Each bucket is stored as a file in partition directory*



Default HDFS location:

`hdfs:///user/hive/warehouse`

Hive Partitioning Model

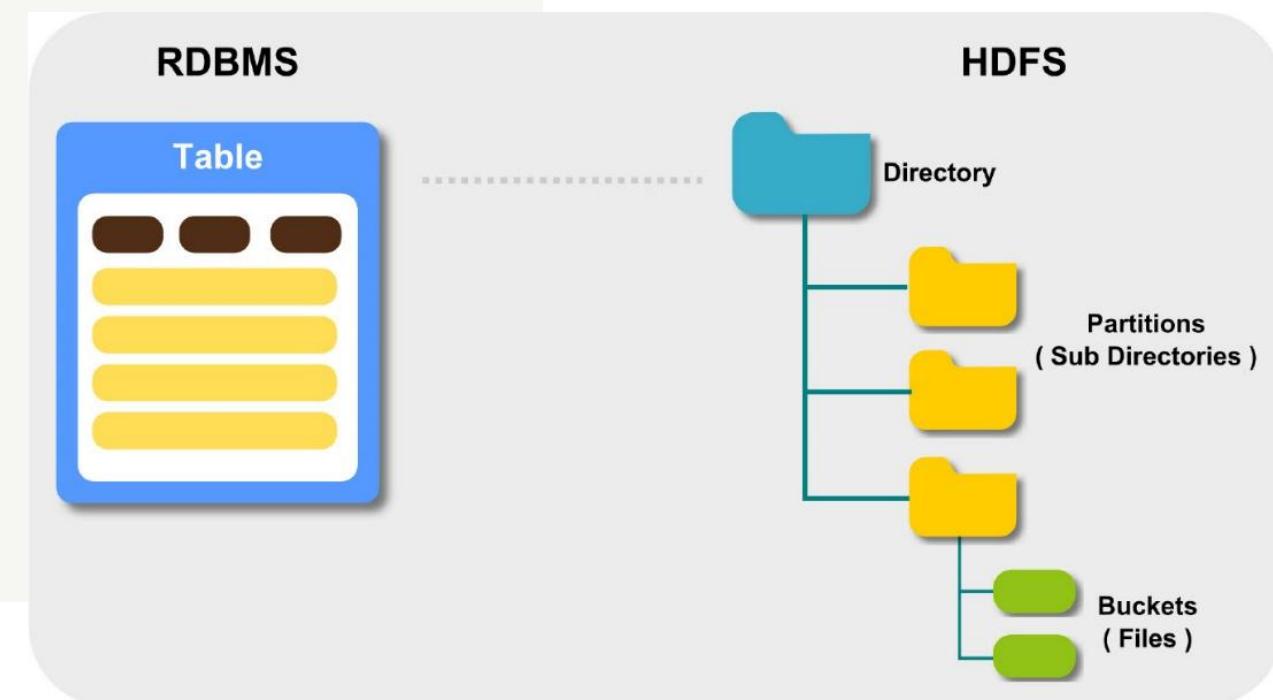


Key idea – Organizing data in a directory tree

- Filter by directories as columns

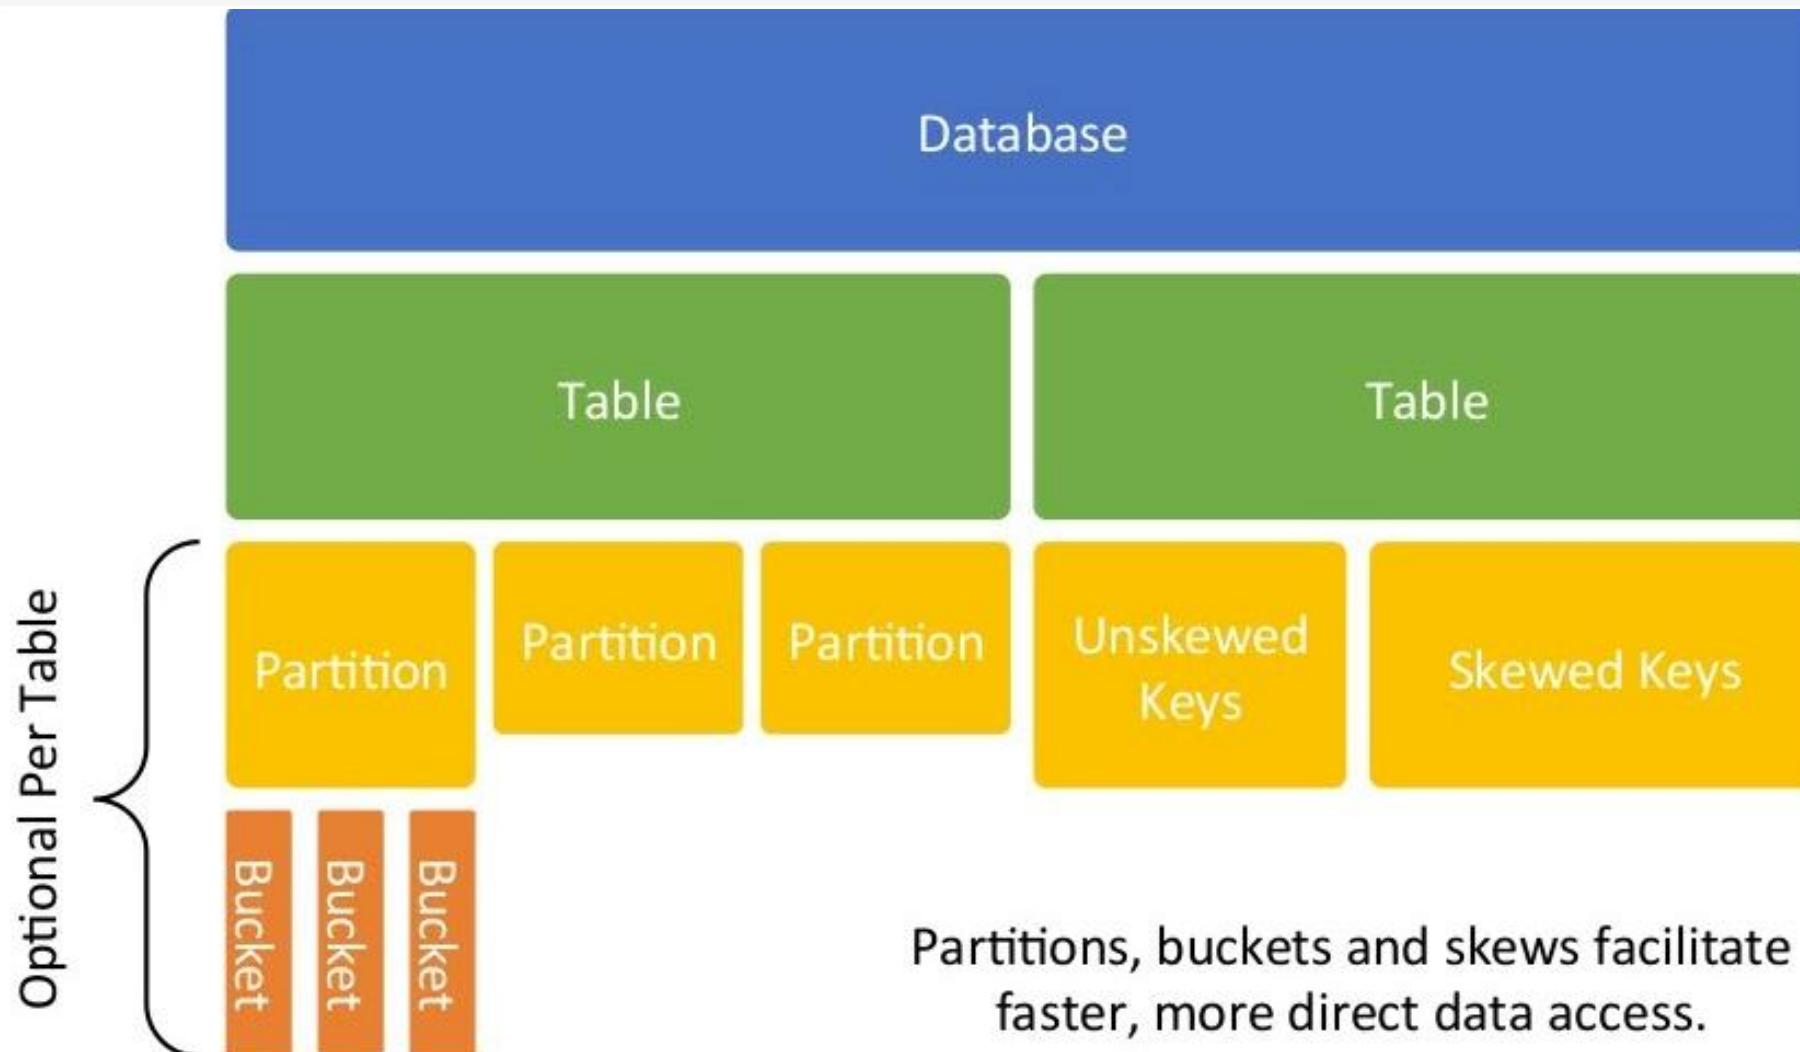
```
SELECT ... WHERE date = '20180513' AND hour = 19
```

```
date=20180513/
|- hour=18/
|  |- ...
|- hour=19/
|  |- 000000_0
|  |- ...
|  |- 000031_0
|- hour=20/
|  |- ...
|- ...
```





Data Abstraction in Hive





Hive Tables Types



Tables

Internal
or
Hive
Managed

*Table's schema and
data are managed by
Hive*

External
or
User
Managed

*Table's schema is
managed by Hive*

*Table's data is
managed by User*

Hive Tables Types



Tables

Drop Internal Tables

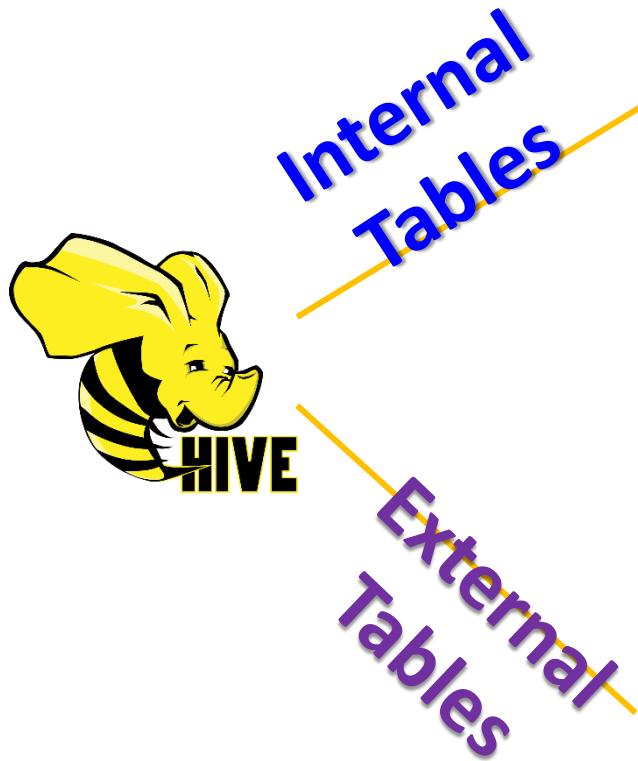
**Will Delete
Schema + Data**

Drop External Tables

**Will Delete
Schema Only**



Hive Tables Types



- **Internal Table (Hive Managed)**

Table Schema and Data are managed by Hive

- We use **Internal** table when:

- Data is temporary
- Data is not needed after table deletion
- Hive **only** is using the table data
(no any external sources like pig, sqoop, mapreduce, ... use the table)

- **External Table (User Managed)**

Table Schema is managed by Hive, Data by the User

- We use **External** table when:

- When data is available in HDFS
- When files are being used **outside** of Hive

HIVE Managed Tables (*Internal*)



1. Create the table and its schema

```
create table products (id int, name string, price float, category string);
```

2. Need to Load the data using the Load / Insert into commands

```
load data inpath /tmp/products.tsv into table products;
```

* products.tsv file is stored on HDFS

- Both data and schema will be removed if the table is dropped

USER Managed Tables (*External*)



A. The data is **already available** on HDFS

1. Create the table and its schema
2. Provide a location on HDFS (directory)

```
create external table products (id int, name string, price float, category string)
location '/tmp/';
```

* *products.tsv* file is stored on HDFS in the /tmp/ directory

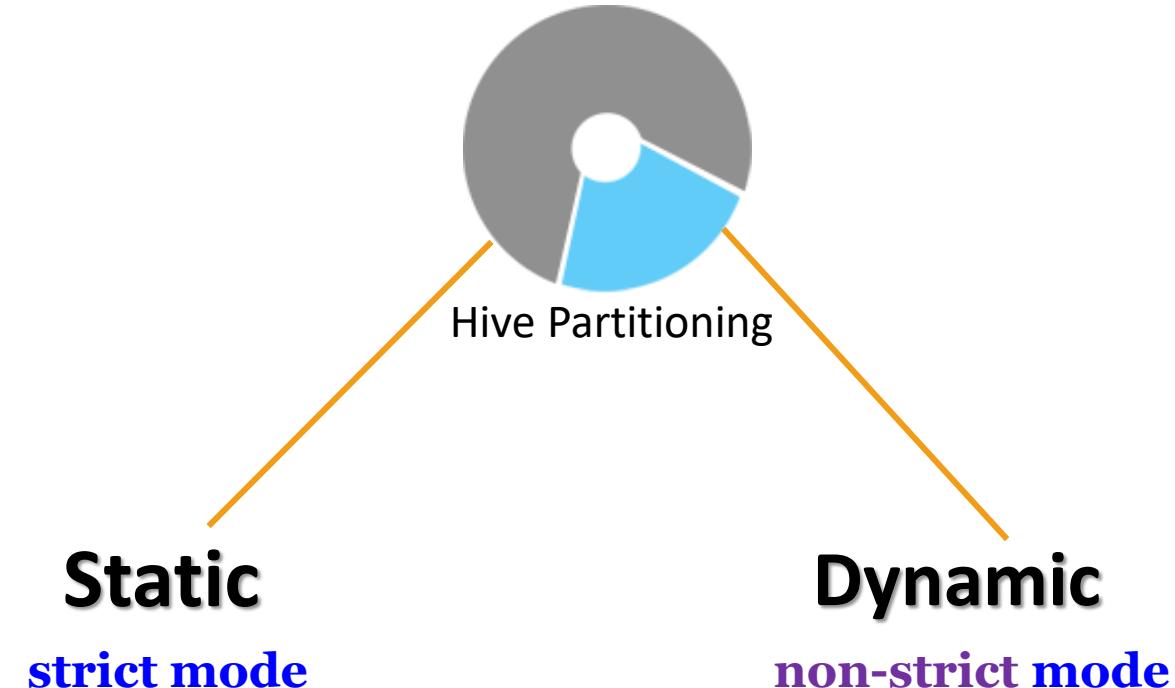
B. Do NOT need a ‘LOAD’ command to load data

● Only schema will be removed if the table is dropped

Hive Table Partitioning



- Can be performed on Hive Managed and Non-Managed tables
- Static partitions** are faster than **dynamic partitions**
- Strict mode** is a useful option in Hive to improve performance and avoid any unexpected delays due to full table scan.



```
set hive.exec.dynamic.partition = true/false;
set hive.exec.dynamic.partition.mode = nonstrict/strict;
```

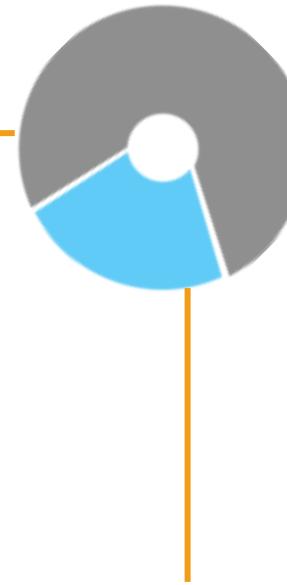
Hive Table Partitioning



● Static Partitioning

- Are preferred when loading big files into Hive tables
- Static partition can be altered
- You “statically” add a partition in the table and move the file into the partition of the table
- You should use where clause to use limit in the static partition.
- Partition column value can be obtained from the filename, day of date etc without reading the whole big file

Hive Partitioning



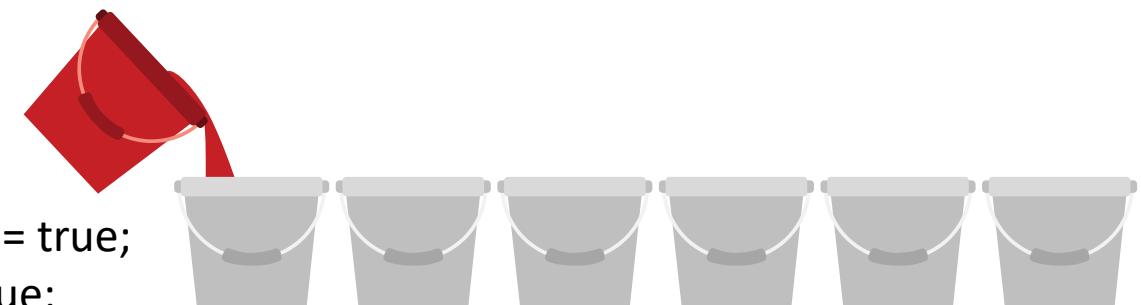
● Dynamic Partitioning

- Single insert to partition table is known as a dynamic partition
- Is used loads the data from the non-partitioned table
- Dynamic partition cannot be altered
- Where clause is not required to use limit

Hive Table Bucketing

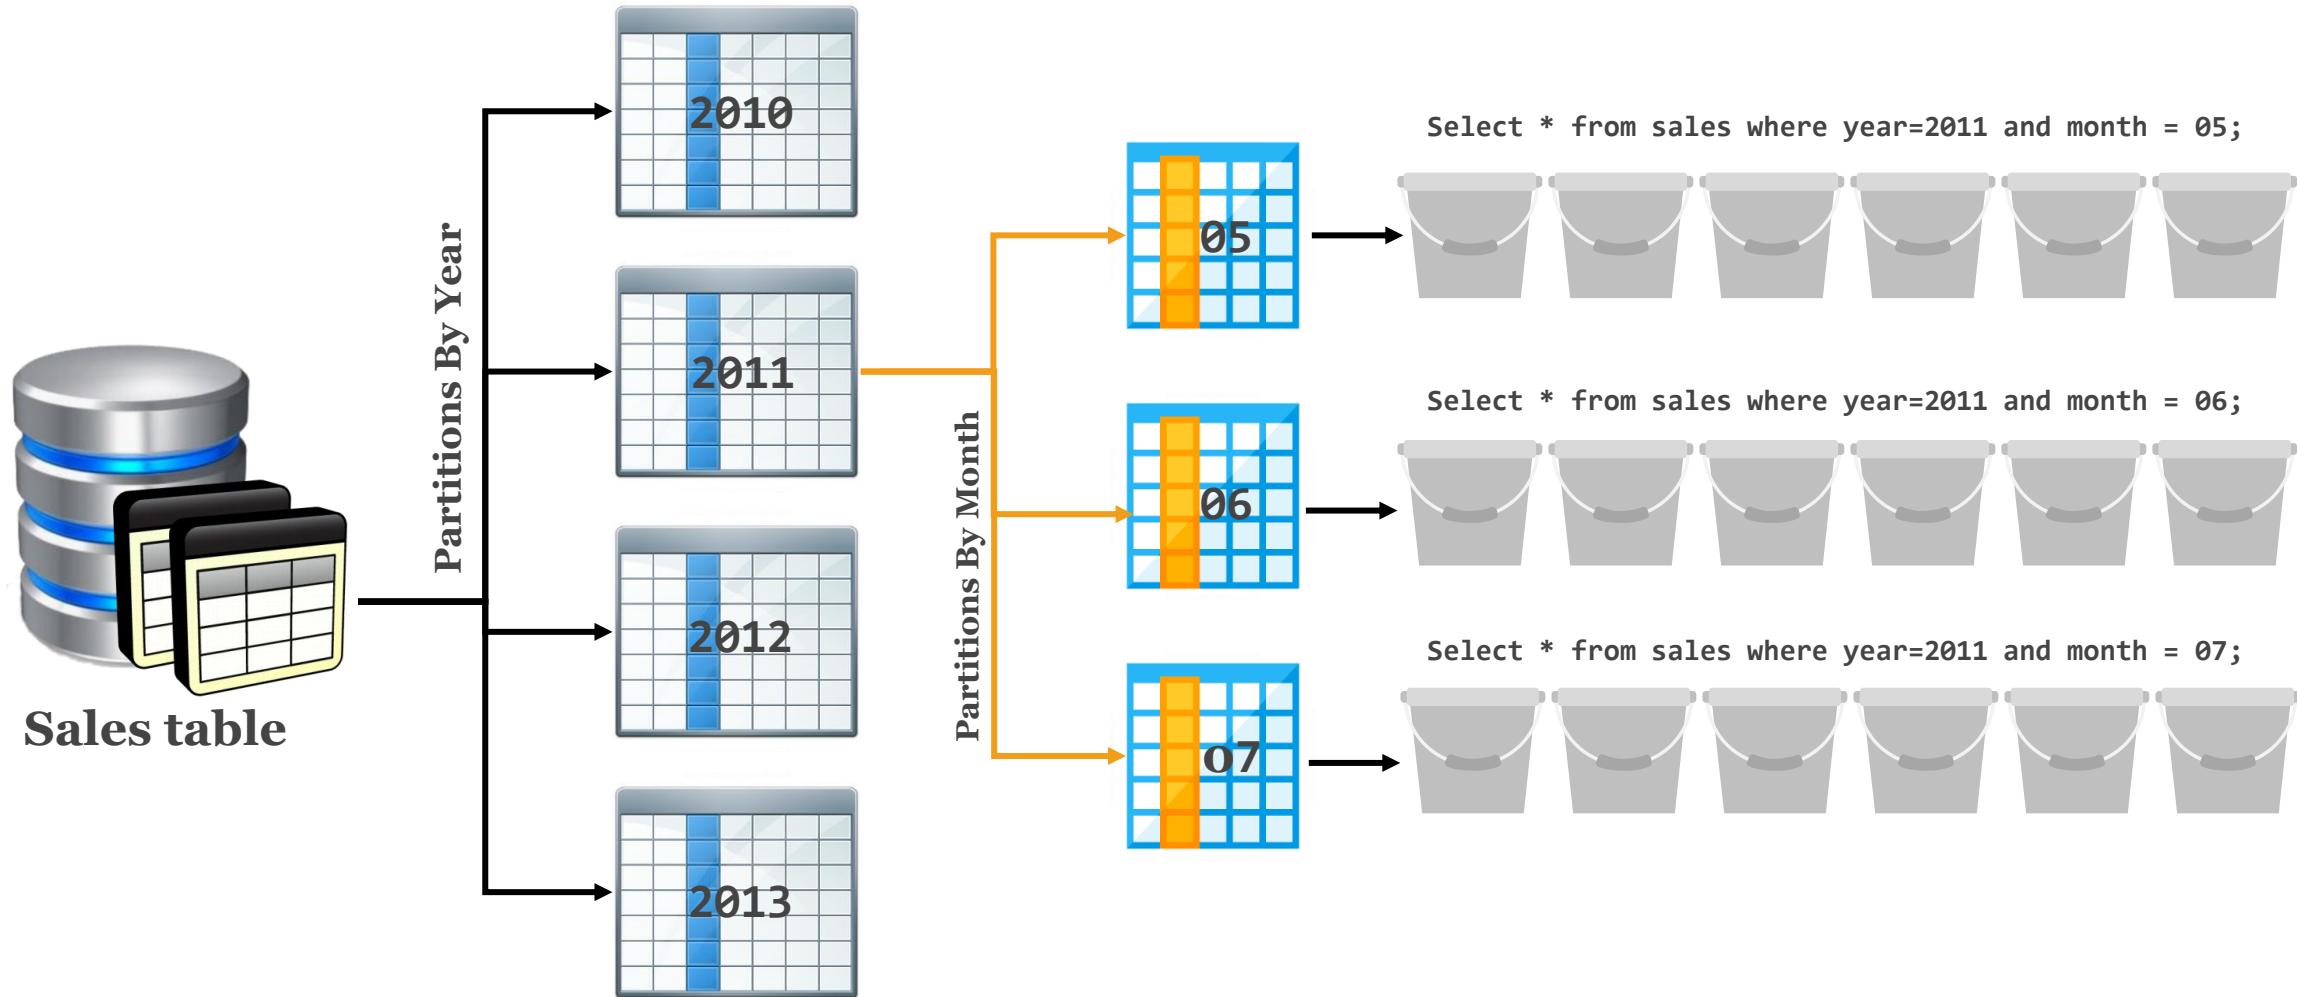


- A technique for decomposing table datasets into more manageable parts
- Based on hashing function on the bucketed column
 - Along with mod (by the total number of buckets)*
- Create almost equally distributed data file parts
- To divide the table into buckets, we use **CLUSTERED BY** clause
- Each bucket is just a file in the table directory. Bucket numbering is 1-based
- Can be done w/without partitioning



```
set hive.enforce.bucketing = true;  
set hive.enforce.sorting=true;
```

Hive Table Partitioning Example



Hive Persistence Formats



● Built-in Formats:

- *ORC File, RCFile* (*Optimized Row Columnar, Row Columnar*)
- *Delimited Text*
- *Avro*
- *Parquet*
- *Sequence File*
- ...



● 3rd-Party Addons (need custom Serde):

- *JSON*
- *XML*
- ...



Loading Data in Hive

- **Hive LOAD statement** *(Not needed for External Hive Tables)*

- *Load files from HDFS or local file system.*
- *Format must agree with table format.*

- **Insert from query**

- *CREATE TABLE AS SELECT or INSERT INTO.*

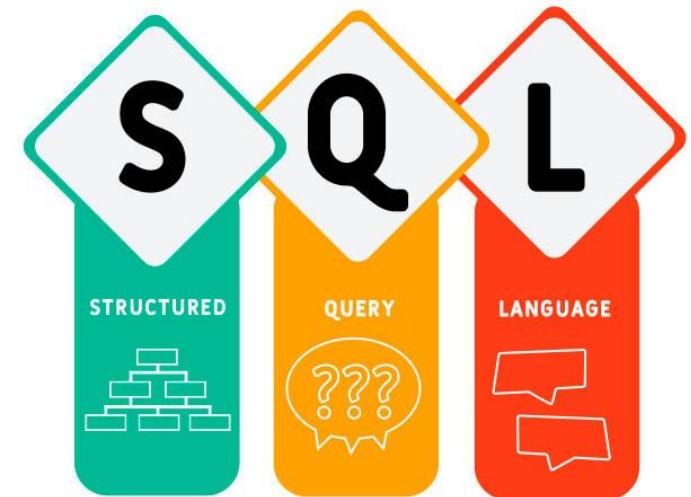
- **Spark / Nifi / ...**

- *Data transfer from external sources to HDFS / Hive.*

Hive Query Language (HiveQL)

- HiveQL / HQL provides the basic SQL-like operations:

- Select columns using **SELECT**
- Filter rows using **WHERE**
- **JOIN** between tables
- Evaluate aggregates using **GROUP BY**
- Store query results into another table
- Download results to a local directory (i.e., export from HDFS)
- Manage tables and queries with **CREATE**, **DROP**, and **ALTER**



Primitive Data Types

Type	Comments
TINYINT, SMALLINT, INT, BIGINT	1, 2, 4 and 8-byte integers
BOOLEAN	TRUE/FALSE
FLOAT, DOUBLE	Single and double precision real numbers
STRING	Character string
TIMESTAMP	Unix-epoch offset or datetime string
DECIMAL	Arbitrary-precision decimal
BINARY	Opaque; ignore these bytes

Complex Data Types

Type	Comments
STRUCT	A collection of elements If S is of type STRUCT {a INT, b INT}: S.a returns element a
MAP	Key-value tuple If M is a map from 'group' to GID: M['group'] returns value of GID
ARRAY	Indexed list If A is an array of elements ['a','b','c']: A[0] returns 'a'

HiveQL : Examples

1/2



- `hive> show tables;`
- `hive> create table shakespeare (freq int, word string)`
`row format delimited fields terminated by '\t' stored as textfile;`
- `hive> load data inpath "shakespeare_freq" into table shakespeare;`
- `hive> select * from shakespeare where freq >100 sort by freq asc limit 10;`

HiveQL : Create Table Examples

2/2



Hive Managed Table

```
create table products (
    id int,
    name string,
    price float,
    category string)
```

row format

delimited fields

terminated by ',';

User Managed Table

```
create EXTERNAL table IF NOT EXISTS Cars (
    name STRING,
    miles_per_Gallon INT,
    cylinders INT,
    horsepower INT,
    Weight_in_lbs INT,
    acceleration DECIMAL,
    `year` DATE)
```

COMMENT 'Data about cars from a public database'

ROW FORMAT DELIMITED FIELDS

TERMINATED BY ',' STORED AS TEXTFILE

Location '/user/username/visdata';

HiveQL : Load / Export Data



Load data into a partitioned table

```
LOAD DATA INPATH '/tmp/pv_2008-06-08_us.txt' INTO TABLE  
page_view PARTITION(date='2008-06-08', country='US');
```

Export data to a file

```
INSERT OVERWRITE DIRECTORY '/tmp/pv_gender_sum'  
SELECT * FROM pv_gender_sum;
```

Hive SerDes



- SerDe is short for **Serializer/Deserializer**
- Hive uses the SerDe interface for IO
- The interface handles both serialization and deserialization and interprets the results of serialization as individual fields for processing
- A SerDe allows Hive to read in data from a table, and write it back out to HDFS in any custom format
- Built-in and Custom SerDes

Anyone can write their own SerDe for their own data formats

Connecting to Hive Server



- **Remotely**

- Zeppelin interface
- beeline (via JDBC, HiveServer2 support)
beeline -u jdbc:hive2://localhost:10000

- JDBC and ODBC drivers (e.g *Squirrel*, *Tableau*, *MS Excel*...)
- Thrift API

- **Run a Hive script**

```
$ hive -f <path>/mysql.hql
```

It's time for a break

Grab some coffee, We'll be back in 15min



Hive Workshop



McGill

School of
Continuing Studies



Hive Query Language

- **Select**
- **FROM**
- **WHERE**
- **GROUP BY**
- **HAVING**
- **JOIN**
- ...

- **Create table**
- **Create view**
- **Drop table**
- **Load**
- ...

Primitive types:

- INTEGERS
 - TINY INT 1 byte integer
 - SMALL INT 2 byte integer
 - INT 4 byte integer
 - BIGINT 8 byte integer
- Date Type
- BOOLEAN
 - BOOLEAN TRUE or FALSE
- FLOATING POINT numbers
 - FLOAT Single precision
 - DOUBLE Double precision
- STRING type
 - STRING Sequence of characters
- . . .

Complex Types: Complex types can be constructed using primitive data types:

- **Structs**
- **Maps** or key value pairs
- **Arrays** – Indexed lists

Create Table Syntax

General Syntax:

```
CREATE TABLE table_name (
    col1 data_type,
    col2 data_type,
    col3 data_type,
    col4 data_type)
```

ROW FORMAT DELIMITED FIELDS

TERMINATED BY '**<char>**'

STORED AS **format_type**;

Example:

```
CREATE TABLE page_view (
    viewTime INT,
    userid BIGINT,
    page_url STRING,
    referrer_url STRING,
    ip STRING COMMENT 'IP Address of the User' )
ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t'
STORED AS TEXTFILE;
```

More Complex Table Example

```
CREATE TABLE employees (
    name STRING,
    salary FLOAT,
    subordinates ARRAY<STRING>,
    deductions MAP<STRING, FLOAT>,
    address STRUCT<street:STRING,
                city:STRING,
                state:STRING,
                zip:INT>)
ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t'
STORED AS TEXTFILE;
```

External Table

```
CREATE EXTERNAL TABLE page_view_stg (
    viewTime INT,
    userid BIGINT,
    page_url STRING,
    referrer_url STRING,
    ip STRING COMMENT 'IP Address of the User')
ROW FORMAT DELIMITED FIELDS
TERMINATED BY '\t'
STORED AS TEXTFILE
LOCATION '/user/staging/page_view';
```

More About Tables

● CREATE TABLE

- **LOAD:** data file **moved** into Hive's data warehouse directory
- **DROP:** **both** metadata and data deleted

● CREATE EXTERNAL TABLE

- **No LOAD and no files moved**
- **DROP:** **only** metadata deleted
- **Use this when sharing with other Hadoop applications, or when you want to use multiple schemas on the same data**

More about Partitioning

- Can make some queries faster
- Divide data based on partition column
- Use **PARTITION BY** clause when creating table
- Use **PARTITION** clause when loading data
- Use **MSCK** to add missing partitions
- **SHOW PARTITIONS** will show a table's partitions

More about Bucketing

- Can speed up queries that involve sampling the data

Sampling allows users to take a subset of dataset and analyze it, without having to analyze the entire data set. If a representative sample is used, then a query can return meaningful results as well as finish quicker and consume fewer compute resources.

- Sampling works without bucketing, but Hive has to scan the entire dataset
- Use **CLUSTERED BY** when creating table
- For sorted buckets, add **SORTED BY**
- To query a sample of your data, use **TABLESAMPLE**



HiveQL : Sampling

Hive offers a built-in TABLESAMPLE clause that allows you to sample your tables.

TABLESAMPLE can return

- *Only subsets of buckets (bucket sampling)*
- *HDFS blocks (block sampling)*
- *Only first N records from each input split.*

INSERT OVERWRITE TABLE pv_gender_sum_sample

**SELECT * FROM pv_gender_sum TABLESAMPLE
(BUCKET 3 OUT OF 32);**

Browsing Tables And Partitions

Command	Comments
<code>SHOW TABLES;</code>	Show all the tables in the database
<code>SHOW TABLES 'page.*';</code>	Show tables matching the specification (uses regex syntax)
<code>SHOW PARTITIONS page_view;</code>	Show the partitions of the page_view table
<code>DESCRIBE page_view;</code>	List columns of the table
<code>DESCRIBE EXTENDED page_view;</code>	More information on columns (useful only for debugging)
<code>DESCRIBE page_view PARTITION (ds='2008-10-31');</code>	List information about a partition

Questions?

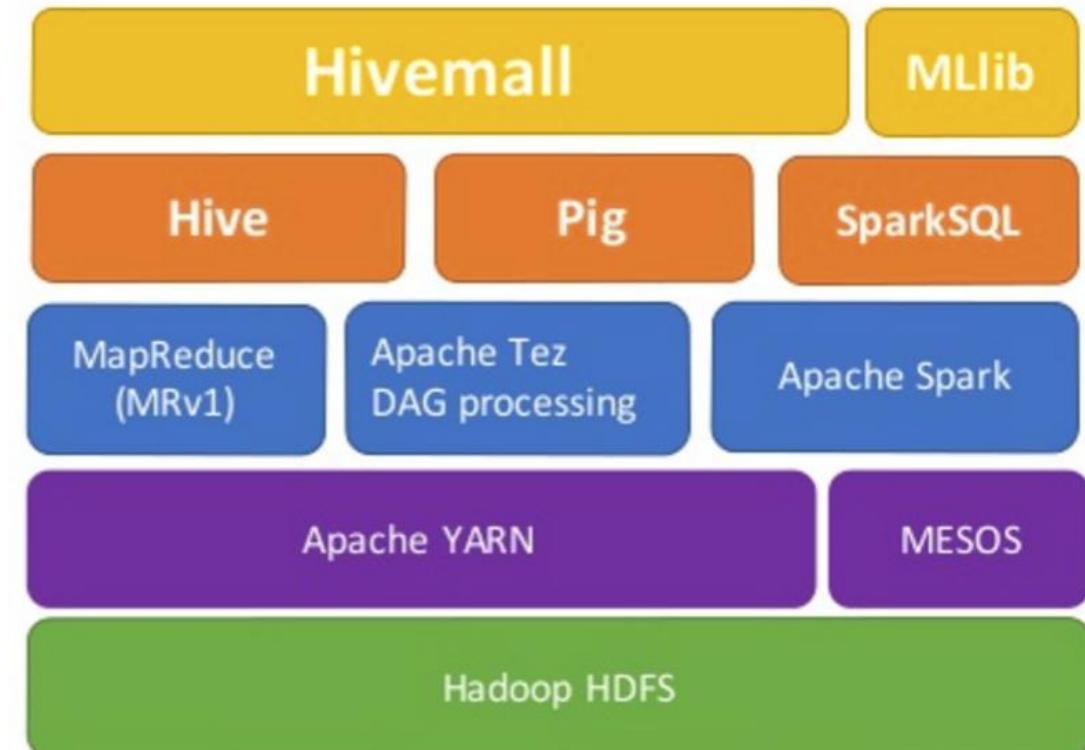


Appendix



Apache HiveMall

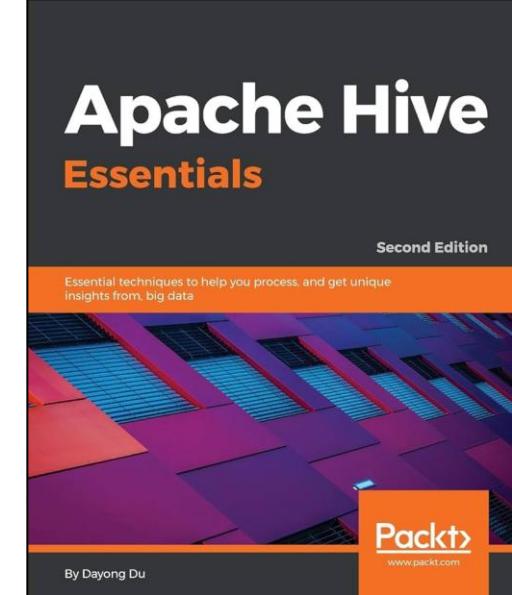
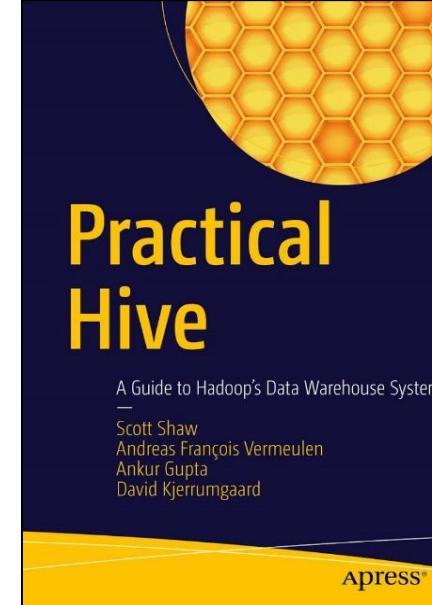
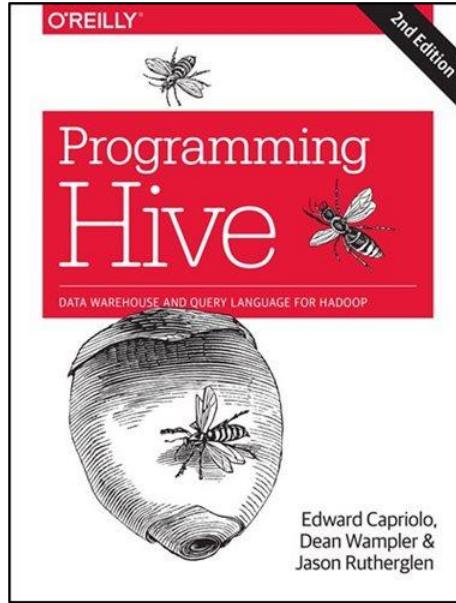
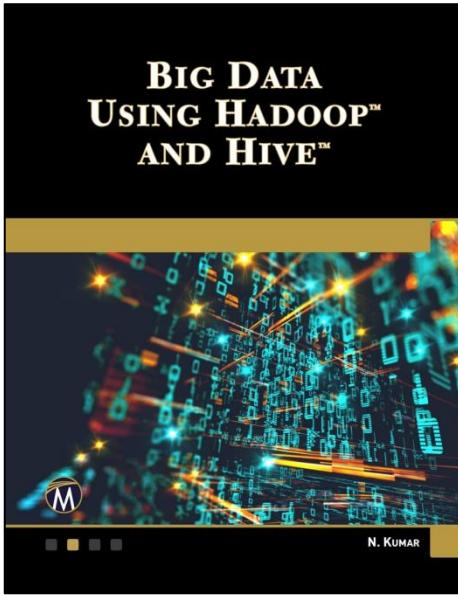
- A scalable machine learning library built as a collections of Hive UDFs.
- Supported Algorithms:
 - Binary Classification
 - Multi-class Classification
 - Regression
 - Recommendation
 - k-Nearest Neighbor
 - Anomaly Detection
 - Natural Language Processing (English/Japanese)



<https://hivemall.incubator.apache.org/>



Resources



Online Resources

<https://hive.apache.org/>

<https://cwiki.apache.org/confluence/display/Hive/LanguageManual>

Merci!
Thank You!

